

Acceleration material

Deep learning

Brice Convers, Jeanne Lemoine

SICOM 3A

19/12/2023



- ① Context
- ② Dataset
- ③ Multiclass classification
- ④ Hyperparameters
- ⑤ Parallelization
- ⑥ Results
- ⑦ Data Augmentation

1 Context

2 Dataset

3 Multiclass classification

4 Hyperparameters

5 Paralellization

6 Results

7 Data Augmentation

Model architecture

UNet architecture




Figure 1: Architecture of a UNet model [2]

Link to the initial model : <https://UNet>

UNet architecture

Pros :

- Specifically designed for semantic segmentation tasks
- Captures context and preserves spatial information
- Performs well even with a limited amount of training data

Cons :

- Sensitive to the input size of images
- Can be computationally intensive

- 1 Context
- 2 Dataset**
- 3 Multiclass classification
- 4 Hyperparameters
- 5 Paralellization
- 6 Results
- 7 Data Augmentation

Data split and resize

Creation of the train and test sets :


- `train_test_split()` function from the *sklearn* library
- Definition of a percentage which separates our dataset into train and test

Resize the images :

- Initial dataset : 400 images (6000x4000 pxls) [5]
- Resize : less computational resources, reduce the memory storage, faster training, reduce overfitting

- 1 Context
- 2 Dataset
- 3 Multiclass classification**
- 4 Hyperparameters
- 5 Paralellization
- 6 Results
- 7 Data Augmentation

Tree of the project



👉 Object-oriented programming with classes

👉 1 file gathering all the parameters
+ 1 main for both test and train

Model architecture

- **Class Block** : store the convolutional and relu layers
- **Class Encoder** : 64 Blocks and maxpool
- **Class Decoder** : up-samplings, cropping and 64 Blocks
- **Class UNet** : assembly of Encoder and Decoder with an interpolation of type `torch.nn.functional.interpolate()`

Adjustments for a binary classification

When $NBR_CLASSES = 1$, we face a binary classification.

👉 Unlabelled binary segmentation

- **Mean** threshold on the masks in grayscale to have 50% of both classes : **extract the contours**




Figure 2: Results for binary classification

- 1 Context
- 2 Dataset
- 3 Multiclass classification
- 4 Hyperparameters**
- 5 Paralellization
- 6 Results
- 7 Data Augmentation

Choice of hyperparameters

- Test split = 0.15 (%) = 340 + 60 images : **not many images so more important training set**
- Batch size = 4, 8 or 16 : **multiple of 4 for 4 GPU**

Smaller batch : slower, less stable, faster convergence, better for GPUs memory, less overfitting

Bigger batch : better generalization

- Model size = 64 : **not too heavy**
- Optimizer = Adam (Adaptive Moment Estimation) : **robust and can adapt the learning rate**
- Loss function = BCEWithlogitsloss() for binary classification or CrossEntropyLoss() for 24 classes

Choice of hyperparameters

- Number of classes = 24
- Image Sizes = (128,128), (256,256), (512,512) : **decrease computational time, small dataset**

For a model with 64 layers we only want the contours, wa can decrease the image size.

- Number of epoch = 10-15 for small test, maximum 50
- Learning rate = 0.01 : **trade-off between over and under fitting and convergence and divergence**

Too big : can oscillate or diverge and miss the solution but faster.

- No dropout or L2 regularization

- 1 Context
- 2 Dataset
- 3 Multiclass classification
- 4 Hyperparameters
- 5 Paralellization**
- 6 Results
- 7 Data Augmentation

Parallelization [3]

- Activation of the parallelism : *ACTIVATE_PARALLELISM = True*
- Generic function for training and testing
- In general, 3 or 4 GPUs with 1 node :
 - ✓ **Improvement of 30% of the time of computation with 3 GPUs**
- No parallelism needed for the test

- 1 Context
- 2 Dataset
- 3 Multiclass classification
- 4 Hyperparameters
- 5 Paralellization
- 6 Results**
- 7 Data Augmentation

Multiclass Segmentation results




Figure 3: Example of output (`CrossEntropyLoss()`, learning rate = 0.01, epoch = 50, image size = (512,512), size of the model = 64, batch size = 4, 340 train images)

Loss curves



Figure 4: Training and testing loss curves

Learning Rate and DICE coefficient

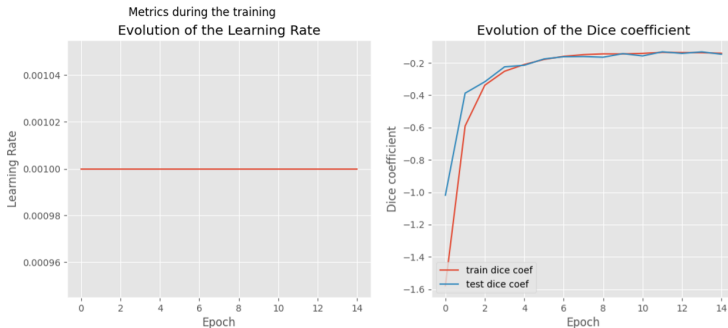


Figure 5: Training Metrics

👉 Convergence of our model

DICE : measure of similarity

Metrics: Confusion Matrix




Figure 6: Confusion Matrix (epoch = 50, 340 training images, learning rate = 0.01)

Metrics: Confusion Matrix

- F1 = 0.0953



Figure 7: Confusion Matrix (epoch = 50, training images, learning rate = 0.01)

- 1 Context
- 2 Dataset
- 3 Multiclass classification
- 4 Hyperparameters
- 5 Paralellization
- 6 Results
- 7 Data Augmentation**

Theory

Why :

- Segmentation was not detailed enough
- Semantic was really bad

How :

- Torch transformations: Horizontal Flip, Vertical Flip, Random Crop
- 400 images and masks → 1600 images and masks [4] [1]

New dataset : 1360 train + 240 test images

Batch size = 4

Metrics: Confusion Matrix

- F1: 0.1023 (up to 7.2%)




Figure 8: Confusion Matrix (epoch = 8, 1360 train images, learning rate = 0.01)

Metrics: Confusion Matrix with AugmentedData

- F1: 0.1021

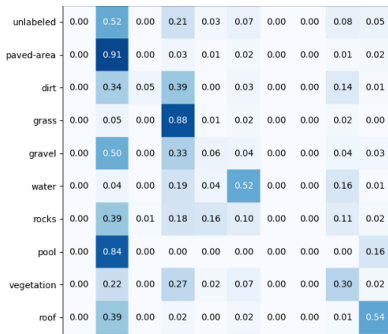


Figure 9: Confusion Matrix (epoch = 50, 1360 train images, learning rate = 0.01)

Early stopping

- Stop the code when the loss function tends towards an asymptotic for several epochs in a row

EARLY_STOPPING_ACTIVATE = True

PATIENCE = 5

Conclusion

+++

- Object-oriented programming model
- Management of the parameters
- Parallelization
- Data augmentation

To be continued :

- Metrics unfinished
- New tests with heavier hyperparameters needed

- [Ayu] Ayu. URL: <https://github.com/ayushdabra/drone-images-semantic-segmentation>.
- [Bel] Nada Belaidi. *U-Net : le réseau de neurones populaire en Computer Vision*. URL: <https://blent.ai/blog/a/unet-computer-vision>.
- [Cha] Dawood Al Chanti. URL: <https://scholar.google.fr/citations?user=osXo54QAAAAJ&hl=en>.
- [Pao] Niccolo Paolinelli. *Semantic Segmentation Drone Dataset (U-Net)*. URL: <https://www.kaggle.com/code/nicopaolinelli/semantic-segmentation-drone-dataset-u-net>.
- [Siy] Bulent Siyah. *Aerial Semantic Segmentation Drone Dataset*. URL: <https://www.kaggle.com/datasets/bulentsiyah/semantic-drone-dataset>.